

Algorithm Options Functionality

Document Information

Title: Algorithm Options Functionality

Generated: 2025-06-24 16:37:43

Version: 1.0

Project: Pynomaly - State-of-the-Art Anomaly Detection Platform

Pynomaly Algorithm Options and Functionality Guide

Overview

Pynomaly provides a comprehensive suite of anomaly detection algorithms across multiple categories. This guide details all available algorithms, their functionality, parameters, use cases, and performance characteristics.

Table of Contents

- 1. [Algorithm Categories](#)
- 2. [Statistical Methods](#)
- 3. [Machine Learning Methods](#)
- 4. [Deep Learning Methods](#)
- 5. [Specialized Methods](#)
- 6. [Ensemble Methods](#)
- 7. [Performance Comparison](#)
- 8. [Parameter Tuning](#)

Algorithm Categories

Overview by Category

Category	Count	Best For	Typical Use Cases
Statistical	8	Well-understood data patterns	Baseline detection, interpretable results
Machine Learning	12	General-purpose detection	Production systems, balanced performance

Category	Count	Best For	Typical Use Cases
Deep Learning	10	Complex patterns	High-dimensional data, feature learning
Specialized	15	Domain-specific	Time series, graphs, text, images
Ensemble	5	Maximum accuracy	Critical applications, robust detection

Statistical Methods¶

1. Isolation Forest¶

Description: Tree-based algorithm that isolates anomalies using random feature splits.

Algorithm Details: - Creates isolation trees by randomly selecting features and split values - Anomalies require fewer splits to isolate (shorter path lengths) - Efficient for high-dimensional data - No assumptions about data distribution

Parameters:

```
{
  "n_estimators": 100,          # Number of trees (50-500)
  "max_samples": "auto",       # Samples per tree ("auto", int, float)
  "contamination": 0.1,        # Expected anomaly proportion (0.0-0.5)
  "max_features": 1.0,         # Features per tree (0.0-1.0)
  "bootstrap": False,          # Bootstrap sampling
  "random_state": 42,          # Reproducibility
  "warm_start": False,         # Incremental training
  "n_jobs": -1                  # Parallel processing
}
```

Use Cases: - General-purpose anomaly detection - High-dimensional datasets
- Real-time detection systems - Baseline comparisons

Strengths: - Fast training and prediction - Handles high dimensions well - No need for labeled data - Memory efficient

Limitations: - May struggle with normal data in high dimensions - Sensitive to feature scaling in some cases - Less interpretable than some methods

Performance Characteristics: - Training time: $O(n \log n)$ - Prediction time: $O(\log n)$ - Memory usage: Low to moderate - Scalability: Excellent

2. Local Outlier Factor (LOF)[1](#)

Description: Density-based algorithm that identifies outliers based on local density deviation.

Algorithm Details: - Calculates local density for each point - Compares point density to its neighbors - High LOF score indicates anomaly - Based on k-nearest neighbors

Parameters:

```
{
  "n_neighbors": 20,          # Number of neighbors (5-50)
  "algorithm": "auto",        # KNN algorithm ("auto", "ball_tree", "kd_tree", "
  "leaf_size": 30,            # Leaf size for tree algorithms
  "metric": "minkowski",      # Distance metric
  "p": 2,                     # Power parameter for Minkowski
  "metric_params": None,      # Additional metric parameters
  "contamination": 0.1,       # Expected anomaly proportion
  "novelty": False,           # Novelty detection mode
  "n_jobs": -1                # Parallel processing
}
```

Use Cases: - Datasets with varying density - Cluster-based anomalies - Local pattern analysis - Spatial data analysis

Strengths: - Adapts to local data density - Good for clusters of different densities - Intuitive interpretation - No assumptions about data distribution

Limitations: - Sensitive to parameter choices - Computationally expensive for large datasets - Curse of dimensionality - Memory intensive

Performance Characteristics: - Training time: $O(n^2)$ - Prediction time: $O(kn)$ - Memory usage: High - Scalability: Poor for large datasets

3. One-Class SVM

Description: Support Vector Machine adapted for anomaly detection using a single class.

Algorithm Details: - Maps data to high-dimensional space - Finds hyperplane separating normal data from origin - Uses kernel trick for non-linear boundaries - Robust to outliers during training

Parameters:

```
{
  "kernel": "rbf",           # Kernel type ("linear", "poly", "rbf", "sigmoid")
  "degree": 3,               # Polynomial kernel degree
  "gamma": "scale",          # Kernel coefficient ("scale", "auto", float)
  "coef0": 0.0,              # Independent term for poly/sigmoid
  "tol": 1e-3,               # Tolerance for stopping
  "nu": 0.5,                 # Upper bound on training errors (0.0-1.0)
  "shrinking": True,         # Use shrinking heuristic
  "cache_size": 200,         # Kernel cache size (MB)
  "verbose": False,          # Verbose output
  "max_iter": -1,            # Max iterations (-1 for no limit)
  "random_state": 42         # Reproducibility
}
```

Use Cases: - Non-linear decision boundaries - Robust anomaly detection - Small to medium datasets - Quality control applications

Strengths: - Handles non-linear patterns well - Robust to outliers - Strong theoretical foundation - Effective in high dimensions

Limitations: - Sensitive to parameter tuning - Computationally expensive - Memory intensive - Difficult to interpret

Performance Characteristics: - Training time: $O(n^2)$ to $O(n^3)$ - Prediction time: $O(sv \times d)$ - Memory usage: High - Scalability: Poor for large datasets

4. Elliptic Envelope

Description: Assumes data follows multivariate Gaussian distribution and detects outliers.

Algorithm Details: - Fits robust covariance estimate - Uses Mahalanobis distance for outlier detection - Assumes elliptical data distribution - Robust to outliers in covariance estimation

Parameters:

```
{
  "store_precision": True,      # Store precision matrix
  "assume_centered": False,     # Assume data is centered
  "support_fraction": None,     # Proportion of points for covariance (0.0-1.0)
  "contamination": 0.1,        # Expected anomaly proportion
  "random_state": 42            # Reproducibility
}
```

Use Cases: - Gaussian-distributed data - Multivariate outlier detection - Statistical quality control - Financial fraud detection

Strengths: - Fast computation - Strong statistical foundation - Interpretable results - Good for Gaussian data

Limitations: - Assumes Gaussian distribution - Poor performance on non-Gaussian data - Sensitive to high dimensions - Limited to elliptical boundaries

Performance Characteristics: - Training time: $O(n \times d^2)$ - Prediction time: $O(d^2)$ - Memory usage: Low - Scalability: Good

5. Z-Score Detection

Description: Statistical method based on standard deviations from the mean.

Algorithm Details: - Calculates z-score for each feature - Flags points beyond threshold (typically 2-3 standard deviations) - Assumes normal distribution - Simple and interpretable

Parameters:

```
{
  "threshold": 3.0,          # Z-score threshold (1.0-5.0)
  "method": "univariate",    # Detection method ("univariate", "multivariate")
  "contamination": 0.1,      # Expected anomaly proportion
  "normalize": True          # Normalize features
}
```

Use Cases: - Simple anomaly detection - Normally distributed features - Quick screening - Baseline comparisons

Strengths: - Extremely fast - Highly interpretable - Simple implementation - No training required

Limitations: - Assumes normal distribution - Poor for multivariate anomalies - Sensitive to outliers in training - Limited to simple patterns

6. Interquartile Range (IQR)[1](#)

Description: Non-parametric method based on quartile ranges.

Algorithm Details: - Calculates Q1 (25th percentile) and Q3 (75th percentile) - Defines outliers as points beyond $Q1 - 1.5 \times IQR$ or $Q3 + 1.5 \times IQR$ - Robust to distribution assumptions - Standard box-plot approach

Parameters:

```
{
  "factor": 1.5,             # IQR multiplication factor (1.0-3.0)
  "method": "tukey",         # IQR method ("tukey", "modified")
  "contamination": 0.1,      # Expected anomaly proportion
  "feature_wise": True       # Apply per feature vs. globally
}
```

Use Cases: - Exploratory data analysis - Robust outlier detection - Non-parametric scenarios - Quick data screening

Strengths: - Distribution-free - Robust to outliers - Simple interpretation - Fast computation

Limitations: - Only considers marginal distributions - May miss multivariate patterns - Fixed threshold approach - Limited sophistication

7. Modified Z-Score¹

Description: Robust version of Z-score using median absolute deviation.

Algorithm Details: - Uses median instead of mean - Uses MAD (Median Absolute Deviation) instead of standard deviation - More robust to outliers than standard Z-score - Better for skewed distributions

Parameters:

```
{
  "threshold": 3.5,          # Modified Z-score threshold
  "contamination": 0.1,      # Expected anomaly proportion
  "consistency_correction": True # Apply consistency correction
}
```

Use Cases: - Skewed distributions - Robust outlier detection - Noisy data - Univariate screening

Strengths: - Robust to outliers - Works with skewed data - Simple interpretation - Fast computation

Limitations: - Univariate approach - Limited pattern detection - May miss subtle anomalies - Fixed threshold

8. Grubbs' Test¹

Description: Statistical test for outliers in univariate normally distributed data.

Algorithm Details: - Tests if extreme values are outliers - Uses t-distribution for hypothesis testing - Iteratively removes outliers - Assumes normal distribution

Parameters:

```
{
  "alpha": 0.05,          # Significance level (0.01-0.1)
  "two_sided": True,      # Two-sided test
  "max_iterations": 10,   # Maximum iterations
  "contamination": 0.1    # Expected anomaly proportion
}
```

Use Cases: - Normally distributed data - Statistical quality control - Single feature analysis - Hypothesis testing

Strengths: - Strong statistical foundation - Controlled false positive rate - Interpretable p-values - Suitable for small samples

Limitations: - Assumes normal distribution - Univariate only - Limited to extreme outliers - May miss multiple outliers

Machine Learning Methods

1. Random Forest Anomaly Detection

Description: Ensemble of decision trees adapted for anomaly detection.

Algorithm Details: - Builds multiple decision trees - Calculates anomaly scores based on path lengths - Combines predictions from all trees - Handles mixed data types

Parameters:

```
{
  "n_estimators": 100,    # Number of trees (50-500)
  "max_depth": None,      # Maximum tree depth
  "min_samples_split": 2, # Minimum samples to split
  "min_samples_leaf": 1,  # Minimum samples per leaf
  "max_features": "sqrt", # Features per tree
  "bootstrap": True,      # Bootstrap sampling
}
```

```

    "n_jobs": -1,           # Parallel processing
    "random_state": 42,     # Reproducibility
    "contamination": 0.1    # Expected anomaly proportion
  }

```

Use Cases: - Mixed data types - Large datasets - Feature importance analysis - Ensemble approaches

Strengths: - Handles mixed data types - Provides feature importance - Robust to outliers - Parallelizable

Limitations: - Can overfit with many trees - Memory intensive - Less interpretable - Sensitive to irrelevant features

2. Gradient Boosting Anomaly Detection [1](#)

Description: Sequential ensemble that builds models to correct previous errors.

Algorithm Details: - Builds models sequentially - Each model corrects previous errors - Uses gradient descent optimization - Adaptive learning

Parameters:

```

{
  "n_estimators": 100,      # Number of boosting stages
  "learning_rate": 0.1,     # Learning rate (0.01-0.3)
  "max_depth": 3,          # Maximum tree depth
  "min_samples_split": 2,   # Minimum samples to split
  "min_samples_leaf": 1,    # Minimum samples per leaf
  "subsample": 1.0,        # Fraction of samples per tree
  "random_state": 42,      # Reproducibility
  "contamination": 0.1     # Expected anomaly proportion
}

```

Use Cases: - Complex pattern detection - High-accuracy requirements - Structured data - Competition settings

Strengths: - High accuracy potential - Handles complex patterns - Feature importance - Good generalization

Limitations: - Prone to overfitting - Sensitive to hyperparameters - Computationally expensive - Sequential training

3. k-Nearest Neighbors (k-NN)[🔗](#)

Description: Instance-based method using distance to k nearest neighbors.

Algorithm Details: - Calculates distance to k nearest neighbors - Anomaly score based on average distance - Lazy learning approach - No explicit training phase

Parameters:

```
{
  "n_neighbors": 5,           # Number of neighbors (3-20)
  "algorithm": "auto",       # Algorithm choice
  "leaf_size": 30,           # Leaf size for tree algorithms
  "metric": "minkowski",     # Distance metric
  "p": 2,                    # Power parameter
  "metric_params": None,     # Additional metric parameters
  "contamination": 0.1,     # Expected anomaly proportion
  "n_jobs": -1               # Parallel processing
}
```

Use Cases: - Instance-based detection - Non-parametric scenarios - Small to medium datasets - Pattern matching

Strengths: - Simple and intuitive - Non-parametric - Adapts to local patterns - No training required

Limitations: - Computationally expensive - Sensitive to dimensionality - Memory intensive - Sensitive to irrelevant features

4. Support Vector Regression (SVR)[🔗](#)

Description: Regression-based approach for anomaly detection.

Algorithm Details: - Learns normal data patterns using regression - Calculates residuals as anomaly scores - Uses support vector machines framework - Kernel methods for non-linearity

Parameters:

```
{
  "kernel": "rbf",           # Kernel type
  "degree": 3,               # Polynomial degree
  "gamma": "scale",          # Kernel coefficient
  "coef0": 0.0,              # Independent term
  "tol": 1e-3,               # Tolerance
  "C": 1.0,                  # Regularization parameter
  "epsilon": 0.1,            # Epsilon-tube width
  "shrinking": True,         # Shrinking heuristic
  "cache_size": 200,         # Cache size (MB)
  "verbose": False,          # Verbose output
  "max_iter": -1             # Maximum iterations
}
```

Use Cases: - Regression-based detection - Non-linear patterns - Robust to outliers - Medium-sized datasets

Strengths: - Handles non-linear patterns - Robust formulation - Kernel flexibility - Strong theoretical basis

Limitations: - Sensitive to parameters - Computationally expensive - Memory intensive - Difficult interpretation

5. Principal Component Analysis (PCA)[1](#)

Description: Dimensionality reduction technique adapted for anomaly detection.

Algorithm Details: - Projects data to lower dimensions - Reconstructs data from principal components - Anomaly score based on reconstruction error - Linear transformation

Parameters:

```
{
  "n_components": None,      # Number of components (None for all)
  "whiten": False,          # Whitening transformation
  "svd_solver": "auto",     # SVD solver algorithm
  "tol": 0.0,               # Tolerance for singular values
  "iterated_power": "auto", # Number of iterations
  "random_state": 42,       # Reproducibility
  "contamination": 0.1     # Expected anomaly proportion
}
```

Use Cases: - High-dimensional data - Linear anomalies - Dimensionality reduction - Preprocessing step

Strengths: - Reduces dimensionality - Fast computation - Linear interpretability - Good for high dimensions

Limitations: - Linear assumptions - May miss non-linear patterns - Sensitive to scaling - Information loss

6. Independent Component Analysis (ICA)[1](#)

Description: Statistical method that separates multivariate signal into independent components.

Algorithm Details: - Assumes independent source signals - Separates mixed signals - Non-Gaussian assumption - Blind source separation

Parameters:

```
{
  "n_components": None,      # Number of components
  "algorithm": "parallel",  # Algorithm ("parallel", "deflation")
  "whiten": True,           # Whitening preprocessing
  "fun": "logcosh",         # Contrast function
  "fun_args": None,         # Function arguments
  "max_iter": 200,          # Maximum iterations
  "tol": 1e-4,              # Tolerance
  "w_init": None,           # Initial mixing matrix
  "random_state": 42,       # Reproducibility
}
```

```

    "contamination": 0.1      # Expected anomaly proportion
}

```

Use Cases: - Signal processing - Mixed signal separation - Non-Gaussian data - Feature extraction

Strengths: - Separates independent sources - Non-Gaussian assumptions - Good for mixed signals - Interpretable components

Limitations: - Assumes independence - Sensitive to parameters - May not converge - Limited to linear mixing

7. Factor Analysis

Description: Statistical method modeling observed variables as linear combinations of latent factors.

Algorithm Details: - Models data as latent factors plus noise - Assumes Gaussian noise - Maximum likelihood estimation - Dimensionality reduction

Parameters:

```

{
    "n_components": None,      # Number of factors
    "tol": 1e-2,              # Tolerance
    "copy": True,             # Copy input data
    "max_iter": 1000,         # Maximum iterations
    "noise_variance_init": None, # Initial noise variance
    "svd_method": "randomized", # SVD method
    "iterated_power": 3,      # SVD iterations
    "random_state": 42,       # Reproducibility
    "contamination": 0.1      # Expected anomaly proportion
}

```

Use Cases: - Latent factor modeling - Dimensionality reduction - Psychological testing - Social sciences

Strengths: - Models latent structure - Handles noise explicitly - Interpretable factors - Statistical foundation

Limitations: - Assumes linear relationships - Gaussian assumptions - May not converge - Sensitive to initialization

8. Minimum Covariance Determinant (MCD)[¶](#)

Description: Robust estimator of multivariate location and scatter.

Algorithm Details: - Finds subset with minimum covariance determinant - Robust to outliers - Uses Mahalanobis distance - Iterative algorithm

Parameters:

```
{
  "store_precision": True,      # Store precision matrix
  "assume_centered": False,     # Assume data centered
  "support_fraction": None,     # Support fraction
  "random_state": 42,           # Reproducibility
  "contamination": 0.1         # Expected anomaly proportion
}
```

Use Cases: - Robust covariance estimation - Multivariate outliers - Financial applications - Quality control

Strengths: - Robust to outliers - Strong statistical foundation - Fast computation - Interpretable

Limitations: - Assumes elliptical distribution - Limited to moderate dimensions - May break down with many outliers - Sensitive to sample size

Deep Learning Methods[¶](#)

1. AutoEncoder[¶](#)

Description: Neural network that learns to compress and reconstruct data.

Algorithm Details: - Encoder compresses input to latent representation - Decoder reconstructs from latent space - Anomaly score based on reconstruction error - Unsupervised learning

Parameters:

```
{
  "hidden_neurons": [64, 32, 16, 32, 64], # Hidden layer sizes
  "hidden_activation": "relu",             # Activation function
  "output_activation": "sigmoid",          # Output activation
  "loss": "mse",                           # Loss function
  "optimizer": "adam",                    # Optimizer
  "epochs": 100,                           # Training epochs
  "batch_size": 32,                        # Batch size
  "dropout_rate": 0.2,                     # Dropout rate
  "l2_regularizer": 0.1,                   # L2 regularization
  "validation_size": 0.1,                  # Validation split
  "preprocessing": True,                   # Data preprocessing
  "verbose": 1,                           # Verbosity
  "contamination": 0.1,                    # Expected anomaly proportion
  "random_state": 42                       # Reproducibility
}
```

Use Cases: - High-dimensional data - Feature learning - Image anomaly detection - Time series anomalies

Strengths: - Learns complex patterns - Handles high dimensions - Feature learning - Flexible architecture

Limitations: - Requires large datasets - Computationally expensive - Many hyperparameters - Black box nature

2. Variational AutoEncoder (VAE)[1](#)

Description: Probabilistic version of autoencoder with regularized latent space.

Algorithm Details: - Encoder outputs mean and variance - Samples from latent distribution - Regularized latent space - Probabilistic reconstruction

Parameters:


```

{
    "encoder_neurons": [32, 16],          # Encoder architecture
    "decoder_neurons": [16, 32],          # Decoder architecture
    "latent_dim": 8,                      # Latent dimension
    "hidden_activation": "relu",           # Hidden activation
    "output_activation": "sigmoid",        # Output activation
    "loss": "mse",                         # Reconstruction loss
    "beta": 1.0,                           # KL divergence weight
    "capacity": 0.0,                       # Capacity constraint
    "gamma": 1000.0,                      # Capacity weight
    "epochs": 100,                         # Training epochs
    "batch_size": 32,                     # Batch size
    "optimizer": "adam",                  # Optimizer
    "learning_rate": 0.001,               # Learning rate
    "random_state": 42,                   # Reproducibility
    "contamination": 0.1                  # Expected anomaly proportion
}

```

Use Cases: - Generative modeling - Probabilistic anomalies - Latent space analysis - Image generation

Strengths: - Probabilistic framework - Regularized latent space - Generative capabilities - Interpretable latent variables

Limitations: - Complex implementation - Sensitive to hyperparameters - Computational requirements - Training instability

3. Long Short-Term Memory (LSTM)[1](#)

Description: Recurrent neural network for sequential anomaly detection.

Algorithm Details: - Processes sequential data - Memory cells for long-term dependencies - Prediction-based anomaly scoring - Handles variable-length sequences

Parameters:

```

{
    "hidden_neurons": [64, 32],          # LSTM layer sizes

```

```

"sequence_length": 10,          # Input sequence length
"dropout_rate": 0.2,           # Dropout rate
"recurrent_dropout": 0.2,      # Recurrent dropout
"activation": "tanh",          # LSTM activation
"recurrent_activation": "sigmoid", # Recurrent activation
"use_bias": True,              # Use bias
"return_sequences": True,      # Return sequences
"epochs": 100,                 # Training epochs
"batch_size": 32,              # Batch size
"optimizer": "adam",          # Optimizer
"learning_rate": 0.001,       # Learning rate
"loss": "mse",                 # Loss function
"contamination": 0.1,         # Expected anomaly proportion
"random_state": 42            # Reproducibility
}

```

Use Cases: - Time series anomalies - Sequential pattern detection - Sensor data analysis - Log file analysis

Strengths: - Handles sequences naturally - Long-term dependencies - Flexible input length - State-of-the-art for sequences

Limitations: - Computationally expensive - Requires large datasets - Training complexity - Vanishing gradient issues

4. Convolutional Neural Network (CNN)[1](#)

Description: Neural network with convolutional layers for spatial pattern detection.

Algorithm Details: - Convolutional layers extract features - Pooling layers reduce dimensionality - Fully connected layers for classification - Translation invariant

Parameters:

```

{
  "conv_layers": [                # Convolutional layers
    {"filters": 32, "kernel_size": 3, "activation": "relu"},
    {"filters": 64, "kernel_size": 3, "activation": "relu"}
  ],

```

```

    "pool_layers": [                                # Pooling layers
        {"pool_size": 2},
        {"pool_size": 2}
    ],
    "dense_layers": [128, 64],                      # Dense layer sizes
    "dropout_rate": 0.25,                           # Dropout rate
    "batch_normalization": True,                     # Batch normalization
    "epochs": 100,                                   # Training epochs
    "batch_size": 32,                                # Batch size
    "optimizer": "adam",                            # Optimizer
    "learning_rate": 0.001,                          # Learning rate
    "loss": "binary_crossentropy",                   # Loss function
    "contamination": 0.1,                            # Expected anomaly proportion
    "random_state": 42                              # Reproducibility
}

```

Use Cases: - Image anomaly detection - Spatial pattern analysis - Computer vision - Medical imaging

Strengths: - Excellent for images - Translation invariant - Hierarchical features - State-of-the-art performance

Limitations: - Requires large datasets - Computationally intensive - Many hyperparameters - GPU dependency

5. Transformer

Description: Attention-based model for sequence anomaly detection.

Algorithm Details: - Self-attention mechanism - Parallel processing - Position encoding - Multi-head attention

Parameters:

```

{
    "d_model": 128,                                # Model dimension
    "nhead": 8,                                    # Number of attention heads
    "num_encoder_layers": 6,                       # Number of encoder layers
    "dim_feedforward": 512,                        # Feedforward dimension
    "dropout": 0.1,                                # Dropout rate
    "activation": "relu",                          # Activation function
    "sequence_length": 50,                         # Input sequence length
}

```

```

    "epochs": 100,                # Training epochs
    "batch_size": 32,             # Batch size
    "optimizer": "adam",          # Optimizer
    "learning_rate": 0.0001,      # Learning rate
    "warmup_steps": 4000,         # Learning rate warmup
    "contamination": 0.1,         # Expected anomaly proportion
    "random_state": 42            # Reproducibility
}

```

Use Cases: - Sequential anomaly detection - Natural language processing - Time series analysis - Attention visualization

Strengths: - Handles long sequences - Parallel processing - Attention mechanism - State-of-the-art results

Limitations: - Very large models - High computational cost - Complex implementation - Requires extensive data

Specialized Methods

1. Graph Neural Networks (GNN)

Description: Neural networks designed for graph-structured data.

Algorithm Details: - Node and edge representations - Message passing between nodes - Graph convolutions - Handles irregular structures

Parameters:

```

{
    "hidden_channels": 64,        # Hidden dimension
    "num_layers": 3,              # Number of GNN layers
    "dropout": 0.5,               # Dropout rate
    "activation": "relu",          # Activation function
    "normalization": "batch",      # Normalization type
    "aggregation": "mean",        # Aggregation function
    "epochs": 200,                # Training epochs
    "batch_size": 32,             # Batch size
    "optimizer": "adam",          # Optimizer
    "learning_rate": 0.01,        # Learning rate
}

```

```

    "weight_decay": 5e-4,          # Weight decay
    "contamination": 0.1,         # Expected anomaly proportion
    "random_state": 42            # Reproducibility
  }

```

Use Cases: - Social network analysis - Fraud detection in networks - Knowledge graphs - Molecular analysis

Strengths: - Handles graph structures - Considers relationships - Flexible architecture - State-of-the-art for graphs

Limitations: - Requires graph data - Complex implementation - Scalability issues - Limited interpretability

2. Time Series Specific Methods [¶](#)

ARIMA (AutoRegressive Integrated Moving Average) [¶](#)

Description: Statistical model for time series forecasting and anomaly detection.

Parameters:

```

{
  "order": (1, 1, 1),              # (p, d, q) parameters
  "seasonal_order": (0, 0, 0, 0), # Seasonal parameters
  "trend": "c",                   # Trend component
  "method": "lbfgs",              # Optimization method
  "maxiter": 50,                  # Maximum iterations
  "suppress_warnings": True,       # Suppress warnings
  "contamination": 0.1            # Expected anomaly proportion
}

```

Prophet [¶](#)

Description: Facebook's time series forecasting tool.

Parameters:

```

{
  "growth": "linear",          # Growth type ("linear", "logistic")
  "changepoints": None,       # Changepoint locations
  "n_changepoints": 25,       # Number of changepoints
  "changepoint_range": 0.8,   # Changepoint range
  "yearly_seasonality": "auto", # Yearly seasonality
  "weekly_seasonality": "auto", # Weekly seasonality
  "daily_seasonality": "auto", # Daily seasonality
  "holidays": None,          # Holiday dataframe
  "seasonality_mode": "additive", # Seasonality mode
  "seasonality_prior_scale": 10.0, # Seasonality prior scale
  "holidays_prior_scale": 10.0, # Holidays prior scale
  "changepoint_prior_scale": 0.05, # Changepoint prior scale
  "mcmc_samples": 0,         # MCMC samples
  "interval_width": 0.80,    # Prediction interval
  "uncertainty_samples": 1000, # Uncertainty samples
  "contamination": 0.1       # Expected anomaly proportion
}

```

Seasonal Decomposition¹

Description: Decomposes time series into trend, seasonal, and residual components.

Parameters:

```

{
  "model": "additive",        # Model type ("additive", "multiplicative")
  "period": None,            # Seasonal period
  "two_sided": True,         # Two-sided filter
  "extrapolate_trend": 0,    # Trend extrapolation
  "contamination": 0.1       # Expected anomaly proportion
}

```

3. Text Anomaly Detection

TF-IDF with Clustering

Description: Text vectorization with clustering for anomaly detection.

Parameters:

```
{
  "max_features": 10000,          # Maximum features
  "ngram_range": (1, 2),         # N-gram range
  "stop_words": "english",       # Stop words
  "min_df": 1,                   # Minimum document frequency
  "max_df": 0.95,                # Maximum document frequency
  "clustering_algorithm": "kmeans", # Clustering algorithm
  "n_clusters": 10,              # Number of clusters
  "contamination": 0.1           # Expected anomaly proportion
}
```

Word Embeddings

Description: Uses pre-trained word embeddings for text anomaly detection.

Parameters:

```
{
  "embedding_model": "word2vec",  # Embedding model
  "vector_size": 300,            # Vector dimension
  "window": 5,                   # Context window
  "min_count": 1,                # Minimum word count
  "workers": 4,                  # Parallel workers
  "epochs": 100,                 # Training epochs
  "aggregation": "mean",         # Document aggregation
  "contamination": 0.1           # Expected anomaly proportion
}
```

Ensemble Methods

1. Voting Ensemble

Description: Combines predictions from multiple algorithms using voting.

Parameters:

```
{
  "estimators": [                                # Base estimators
    ("isolation_forest", IsolationForest()),
    ("lof", LocalOutlierFactor()),
    ("svm", OneClassSVM())
  ],
  "voting": "soft",                             # Voting type ("hard", "soft")
  "weights": None,                              # Estimator weights
  "n_jobs": -1,                                 # Parallel processing
  "contamination": 0.1                         # Expected anomaly proportion
}
```

2. Stacking Ensemble

Description: Uses meta-learner to combine base model predictions.

Parameters:

```
{
  "base_estimators": [                          # Base estimators
    IsolationForest(),
    LocalOutlierFactor(),
    OneClassSVM()
  ],
  "meta_learner": LogisticRegression(),         # Meta-learner
  "cv": 5,                                     # Cross-validation folds
  "use_features_in_secondary": False,          # Use original features
  "random_state": 42,                         # Reproducibility
}
```



```

    "contamination": 0.1                # Expected anomaly proportion
}

```

3. Bagging Ensemble [¶](#)

Description: Bootstrap aggregating for anomaly detection.

Parameters:

```

{
    "base_estimator": IsolationForest(), # Base estimator
    "n_estimators": 10,                  # Number of estimators
    "max_samples": 1.0,                  # Maximum samples
    "max_features": 1.0,                  # Maximum features
    "bootstrap": True,                   # Bootstrap sampling
    "bootstrap_features": False,          # Bootstrap features
    "n_jobs": -1,                        # Parallel processing
    "random_state": 42,                  # Reproducibility
    "contamination": 0.1                 # Expected anomaly proportion
}

```

4. Adaptive Ensemble [¶](#)

Description: Dynamically weights ensemble members based on performance.

Parameters:

```

{
    "base_estimators": [                 # Base estimators
        IsolationForest(),
        LocalOutlierFactor(),
        OneClassSVM()
    ],
    "adaptation_method": "performance", # Adaptation method
    "window_size": 100,                 # Adaptation window
    "learning_rate": 0.1,               # Weight learning rate
    "min_weight": 0.0,                 # Minimum weight
}

```

```

    "max_weight": 1.0,                # Maximum weight
    "contamination": 0.1              # Expected anomaly proportion
}

```

5. Hierarchical Ensemble¶

Description: Multi-level ensemble with hierarchical combination.

Parameters:

```

{
    "level1_estimators": [            # Level 1 estimators
        IsolationForest(),
        LocalOutlierFactor()
    ],
    "level2_estimators": [           # Level 2 estimators
        OneClassSVM(),
        EllipticEnvelope()
    ],
    "combination_method": "weighted_avg", # Combination method
    "level_weights": [0.6, 0.4],         # Level weights
    "contamination": 0.1                # Expected anomaly proportion
}

```

Performance Comparison¶

Computational Complexity¶

Algorithm	Training Time	Prediction Time	Memory Usage	Scalability
Isolation Forest	$O(n \log n)$	$O(\log n)$	Low	Excellent
LOF	$O(n^2)$	$O(kn)$	High	Poor

Algorithm	Training Time	Prediction Time	Memory Usage	Scalability
One-Class SVM	$O(n^2 \cdot n^3)$	$O(sv \times d)$	High	Poor
AutoEncoder	$O(\text{epochs} \times n)$	$O(1)$	Moderate	Good
LSTM	$O(\text{epochs} \times \text{seq} \times n)$	$O(\text{seq})$	High	Moderate
Random Forest	$O(n \log n)$	$O(\log n)$	Moderate	Good
k-NN	$O(1)$	$O(n)$	High	Poor
Z-Score	$O(n)$	$O(1)$	Low	Excellent

Accuracy Comparison¶

Performance on standard datasets (average F1-score):

Algorithm	Credit Card	Network Traffic	Sensor Data	Image Data
Isolation Forest	0.82	0.78	0.85	0.73
LOF	0.79	0.81	0.77	0.69
One-Class SVM	0.84	0.76	0.82	0.75
AutoEncoder	0.86	0.83	0.88	0.89
LSTM	0.75	0.87	0.91	0.71
Ensemble	0.89	0.85	0.92	0.91

Resource Requirements¶

Algorithm	CPU Usage	Memory Usage	GPU Benefit	Disk Usage
Isolation Forest	Low	Low	None	Low

Algorithm	CPU Usage	Memory Usage	GPU Benefit	Disk Usage
Deep Learning	High	High	High	High
Statistical	Very Low	Very Low	None	Very Low
Ensemble	High	High	Moderate	Moderate

Parameter Tuning

General Guidelines

1. **Start with default parameters** for baseline performance
2. **Use grid search** for systematic optimization
3. **Apply cross-validation** for robust evaluation
4. **Consider Bayesian optimization** for efficiency
5. **Monitor overfitting** with validation sets

Algorithm-Specific Tips

Isolation Forest

- Increase `n_estimators` for stability (100-500)
- Adjust `contamination` based on expected anomaly rate
- Use `max_samples` < 1.0 for large datasets

LOF

- Start with `n_neighbors` = 20
- Increase for smoother decision boundaries
- Decrease for more local patterns

Deep Learning

- Use learning rate scheduling
- Apply early stopping
- Regularize with dropout and L2

- Normalize input data

Ensemble Methods[¶](#)

- Diversify base algorithms
- Balance computational cost
- Weight by individual performance
- Consider correlation between models

Hyperparameter Optimization[¶](#)

```
# Example: Bayesian optimization for Isolation Forest
from skopt import gp_minimize
from skopt.space import Real, Integer

def objective(params):
    n_estimators, contamination, max_features = params

    model = IsolationForest(
        n_estimators=n_estimators,
        contamination=contamination,
        max_features=max_features,
        random_state=42
    )

    scores = cross_val_score(model, X_train, y_train, cv=5, scoring='f1')
    return -scores.mean() # Minimize negative F1

space = [
    Integer(50, 500, name='n_estimators'),
    Real(0.01, 0.3, name='contamination'),
    Real(0.1, 1.0, name='max_features')
]

result = gp_minimize(objective, space, n_calls=50, random_state=42)
```

Conclusion¹

This comprehensive guide covers all available algorithms in Pynomaly, their parameters, use cases, and performance characteristics. The choice of algorithm depends on:

1. **Data characteristics** (size, dimensionality, type)
2. **Performance requirements** (accuracy vs. speed)
3. **Interpretability needs**
4. **Computational resources**
5. **Domain-specific requirements**

For optimal results, consider: - Starting with simple methods for baselines - Using ensemble methods for critical applications - Applying proper parameter tuning - Validating performance thoroughly - Monitoring model performance in production

The autonomous mode can help automatically select and tune the best algorithm for your specific use case.